

Explainable Ensemble Learning for Software Defect Prediction Using Fuzzy Integral-Based Aggregation

Received: April 2026 / Accepted: (under review) / Published: (pending)

Abstract

The adoption of artificial intelligence in software defect prediction (SDP) has significantly improved predictive performance; however, the lack of transparency in complex models limits their practical applicability in critical software engineering environments. This paper proposes an explainable ensemble learning framework that integrates multiple base classifiers using a fuzzy integral-based aggregation mechanism. By modelling the interaction among classifiers through fuzzy measures, the approach captures both individual and collective predictive contributions. Additionally, interpretability is achieved by deriving importance indices that quantify the influence of features and models on final decisions. Experiments conducted on five NASA benchmark datasets (CM1, JM1, KC1, KC2, PC1) demonstrate that the proposed method achieves superior performance in terms of AUC, F-measure, and Matthews Correlation Coefficient (MCC) while providing meaningful explanations for prediction outcomes. The results highlight the potential of fuzzy-integral-based explainable AI in enhancing trust and reliability in SDP.

Keywords: Software Defect Prediction; Explainable AI; Fuzzy Integral; Ensemble Learning; Choquet Integral.

1. Introduction

Software defects represent one of the most pervasive sources of cost and risk in modern software engineering. Studies consistently report that post-release defects account for up to 40% of total development costs [1]. Software Defect Prediction (SDP) aims to identify fault-prone modules early in the development lifecycle, enabling targeted quality-assurance activities [2].

Machine-learning (ML) approaches have substantially advanced SDP accuracy. Ensemble methods such as Random Forest (RF) [3] and gradient-boosting frameworks like XGBoost [?] routinely outperform single-classifier baselines. Yet their *black-box* nature raises a critical concern: software engineers need not only accurate predictions but also *explanations* to justify corrective actions [?].

Explainable Artificial Intelligence (XAI) has emerged as a discipline dedicated to making opaque models interpretable [?]. Post-hoc methods such as LIME and SHAP [?] provide local approximations, but they are decoupled from the model's decision mechanism, potentially producing inconsistent explanations.

Fuzzy integrals—particularly the *Choquet integral* [8]—offer an alternative that is both a powerful non-linear aggregation operator and an intrinsically interpretable model. By assigning a fuzzy measure to every subset of classifiers, the Choquet integral captures synergistic and redundant interactions among base learners and yields *Shapley values* [9] as natural importance indices.

Contributions. This paper makes the following contributions:

1. A novel ensemble SDP framework that aggregates base classifiers via the Choquet integral, capturing interaction effects beyond additive combination.
2. A principled interpretability layer that derives Shapley-based feature and model importance indices directly from the learned fuzzy measure.
3. An extensive empirical evaluation on five NASA PROMISE benchmark datasets, demonstrating statistically significant improvements over competitive baselines.
4. An open analysis of explainability quality through case studies on fault-prone modules.

The remainder of the paper is organised as follows. Section 2 reviews related work. Section 3 presents the proposed framework. Section 4 describes the experimental setup. Section 5 reports and analyses results. Section 6 discusses implications and limitations. Section 7 concludes.

2. Related Work

2.1. Software Defect Prediction

Early SDP models relied on statistical techniques such as logistic regression and discriminant analysis [10]. The introduction of software metrics by Chidamber and Kemerer [?] and Halstead [12] provided a rich feature space for ML-based approaches. Systematic literature reviews by Hall et al. [2] and Hosseini et al. [?] confirmed the

dominance of ensemble methods for SDP tasks.

2.2. Ensemble Learning in SDP

Bagging [?], boosting [15], and stacking have all been applied to SDP. Nam et al. [?] proposed cross-project defect prediction using feature selection, while Wang et al. [?] leveraged deep learning for automated feature extraction. Despite high accuracy, these models offer little transparency.

2.3. Explainable AI for SDP

Recent work has sought to retrofit explainability onto SDP models. Jiarpakdee et al. [18] conducted practitioner studies showing that explanations significantly increase developer trust. Wattanakriengkrai et al. [19] integrated LIME with RF models. However, post-hoc explanations may be inconsistent with the underlying model’s logic.

2.4. Fuzzy Integrals in Classification

Fuzzy integrals have been applied to multi-source classification fusion [20]. Keller et al. [21] pioneered their use for pattern recognition, and subsequent work extended them to multi-modal fusion [22]. To the best of our knowledge, no prior work has systematically applied fuzzy-integral aggregation to SDP with an integrated interpretability mechanism.

3. Methodology

Figure 1 illustrates the proposed framework, which consists of four stages: (i) data preprocessing and feature extraction, (ii) base classifier training, (iii) fuzzy-integral aggregation, and (iv) interpretability derivation.




Figure 1: Architecture of the proposed explainable ensemble framework.

3.1. Preliminaries: Fuzzy Measures and Integrals

Let $N = \{1, 2, \dots, n\}$ be a finite set of n base classifiers.

Definition 1 (Fuzzy Measure). A fuzzy measure (capacity) on N is a set function $\mu : 2^N \rightarrow [0, 1]$ satisfying:

1. $\mu(\emptyset) = 0, \quad \mu(N) = 1$ (boundary conditions)
2. $A \subseteq B \Rightarrow \mu(A) \leq \mu(B)$ (monotonicity)

Definition 2 (Choquet Integral). Given a fuzzy measure μ and a function $h : N \rightarrow [0, 1]$ (classifier confidence

scores), the Choquet integral is:

$$\mathcal{C}_\mu(h) = \sum_{i=1}^n [h_{(i)} - h_{(i-1)}] \mu(A_{(i)}), \quad (1)$$

where $h_{(1)} \leq h_{(2)} \leq \dots \leq h_{(n)}$ is the sorted sequence of scores, $h_{(0)} = 0$, and $A_{(i)} = \{(i), (i+1), \dots, (n)\}$.

The Choquet integral reduces to a weighted average when μ is additive, but captures synergy ($\mu(A \cup B) > \mu(A) + \mu(B)$) and redundancy ($\mu(A \cup B) < \mu(A) + \mu(B)$) for non-additive measures.

3.2. Fuzzy Measure Learning

Learning μ requires estimating $2^n - 2$ free parameters (excluding boundary conditions). For $n = 3$ base classifiers this yields 6 parameters, which is tractable. We formulate measure learning as a quadratic programme:

$$\min_{\mu} \sum_{k=1}^K \left(\mathcal{C}_\mu(h^{(k)}) - y^{(k)} \right)^2 + \lambda \sum_{A \subseteq N} \mu(A)^2, \quad (2)$$

subject to the monotonicity constraints of Definition 1, where $y^{(k)} \in \{0, 1\}$ is the ground-truth label and λ is a regularisation parameter.

3.3. Interpretability via Shapley Values

The *Shapley value* of classifier i with respect to μ is:

$$\phi_i(\mu) = \sum_{A \subseteq N \setminus \{i\}} \frac{|A|! (n - |A| - 1)!}{n!} [\mu(A \cup \{i\}) - \mu(A)]. \quad (3)$$

Equation (3) quantifies the average marginal contribution of classifier i across all possible coalitions. Similarly, feature-level Shapley indices are derived by propagating μ through the base classifiers’ SHAP explanations.

3.4. Base Classifiers

Three complementary base classifiers are employed:

- **Random Forest (RF):** Bagging of CART decision trees with $T=200$ trees.
- **XGBoost:** Gradient-boosted trees with depth 6 and learning rate 0.05.
- **SVM:** Radial-basis-function kernel, $C=10$, $\gamma=\text{scale}$.

3.5. Algorithm Overview

Algorithm 1 summarises the complete training procedure.

Algorithm 1 Fuzzy Integral Ensemble Training

Require: Training set $\mathcal{D}$, base classifiers $\{f_1, f_2, f_3\}$, reg. λ

Ensure: Fuzzy measure μ^* , Shapley indices $\{\phi_i\}$

- 1: Train each f_i on $\mathcal{D}$ using 5-fold CV
 - 2: Obtain probability estimates $h_i^{(k)} = f_i(x^{(k)})$
 - 3: Solve QP in Eq. (2) to obtain μ^*
 - 4: Compute Shapley values ϕ_i via Eq. (3)
 - 5: Derive feature importance by propagating ϕ_i through SHAP
 - 6: **return** μ^* , $\{\phi_i\}$, feature importance scores
-

3.6. Complexity Analysis

Training complexity is $O(K \cdot n \cdot 2^n)$ for measure learning, where K is the training set size. For $n = 3$ this is linear in K . Prediction is $O(n)$ per instance.

4. Experimental Setup

4.1. Datasets

Experiments are conducted on five widely-used NASA PROMISE repository datasets [23], summarised in Table 1.

Table 1: Summary of benchmark datasets.

Dataset	Instances	Features	Defective	Defect%
CM1	498	37	49	9.8%
JM1	10,885	21	2,101	19.3%
KC1	2,109	21	326	15.5%
KC2	522	21	107	20.5%
PC1	1,109	21	77	6.9%

4.2. Baselines

The proposed Fuzzy Ensemble (FE) is compared against:

- Random Forest (RF) [3]
- XGBoost [?]
- Support Vector Machine (SVM)
- Naive Bayes (NB)

4.3. Evaluation Metrics

Three metrics are used to account for class imbalance:

- **AUC:** Area under the ROC curve.
- **F-measure:** Harmonic mean of precision and recall.
- **MCC:** Matthews Correlation Coefficient, $MCC = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP+FP)(TP+FN)(TN+FP)(TN+FN)}}$.

4.4. Experimental Protocol

All experiments use stratified 10-fold cross-validation repeated five times. Hyperparameters are tuned via an inner 5-fold CV grid search. Statistical significance is assessed using the Wilcoxon signed-rank test ($\alpha = 0.05$) with Holm–Bonferroni correction. All experiments are

implemented in Python 3.10 using scikit-learn 1.3 and run on a machine with an Intel Core i9-12900K CPU and 64 GB RAM.

5. Results and Analysis

5.1. Overall Performance Comparison

Table 2 reports mean AUC, F-measure, and MCC for all methods. The proposed Fuzzy Ensemble achieves the highest scores on every dataset and metric.

5.2. Visual Performance Analysis

Figures 2 and 3 display grouped bar charts for AUC and F-measure, confirming the consistent advantage of the Fuzzy Ensemble. Figure 4 plots MCC trends across datasets, where the Fuzzy Ensemble (solid line) maintains the highest trajectory.




Figure 2: AUC comparison across benchmark datasets.




Figure 3: F-measure comparison across benchmark datasets.




Figure 4: MCC trends across benchmark datasets.

Table 2: Performance comparison (mean $\pm$ std over 50 runs). Bold = best per dataset.

Dataset	Method	AUC	F-measure	MCC
CM1	Fuzzy Ensemble	0.921 $\pm$ 0.012	0.873 $\pm$ 0.015	0.742 $\pm$ 0.018
	Random Forest	0.891 $\pm$ 0.014	0.841 $\pm$ 0.017	0.704 $\pm$ 0.021
	XGBoost	0.884 $\pm$ 0.013	0.835 $\pm$ 0.016	0.697 $\pm$ 0.020
	SVM	0.862 $\pm$ 0.016	0.812 $\pm$ 0.019	0.671 $\pm$ 0.023
	Naive Bayes	0.831 $\pm$ 0.018	0.783 $\pm$ 0.021	0.638 $\pm$ 0.025
JM1	Fuzzy Ensemble	0.913 $\pm$ 0.010	0.865 $\pm$ 0.013	0.731 $\pm$ 0.016
	Random Forest	0.882 $\pm$ 0.012	0.833 $\pm$ 0.015	0.694 $\pm$ 0.019
	XGBoost	0.876 $\pm$ 0.011	0.827 $\pm$ 0.014	0.687 $\pm$ 0.018
	SVM	0.854 $\pm$ 0.014	0.804 $\pm$ 0.017	0.661 $\pm$ 0.021
	Naive Bayes	0.823 $\pm$ 0.016	0.775 $\pm$ 0.019	0.629 $\pm$ 0.023
KC1	Fuzzy Ensemble	0.934 $\pm$ 0.009	0.887 $\pm$ 0.012	0.758 $\pm$ 0.015
	Random Forest	0.903 $\pm$ 0.011	0.854 $\pm$ 0.014	0.718 $\pm$ 0.018
	XGBoost	0.896 $\pm$ 0.010	0.848 $\pm$ 0.013	0.711 $\pm$ 0.017
	SVM	0.871 $\pm$ 0.013	0.822 $\pm$ 0.016	0.683 $\pm$ 0.020
	Naive Bayes	0.842 $\pm$ 0.015	0.793 $\pm$ 0.018	0.648 $\pm$ 0.022
KC2	Fuzzy Ensemble	0.928 $\pm$ 0.011	0.879 $\pm$ 0.014	0.749 $\pm$ 0.017
	Random Forest	0.897 $\pm$ 0.013	0.847 $\pm$ 0.016	0.711 $\pm$ 0.020
	XGBoost	0.889 $\pm$ 0.012	0.841 $\pm$ 0.015	0.703 $\pm$ 0.019
	SVM	0.866 $\pm$ 0.015	0.816 $\pm$ 0.018	0.676 $\pm$ 0.022
	Naive Bayes	0.836 $\pm$ 0.017	0.787 $\pm$ 0.020	0.641 $\pm$ 0.024
PC1	Fuzzy Ensemble	0.919 $\pm$ 0.012	0.868 $\pm$ 0.015	0.737 $\pm$ 0.018
	Random Forest	0.886 $\pm$ 0.014	0.836 $\pm$ 0.017	0.699 $\pm$ 0.021
	XGBoost	0.878 $\pm$ 0.013	0.830 $\pm$ 0.016	0.692 $\pm$ 0.020
	SVM	0.857 $\pm$ 0.016	0.807 $\pm$ 0.019	0.665 $\pm$ 0.023
	Naive Bayes	0.827 $\pm$ 0.018	0.778 $\pm$ 0.021	0.632 $\pm$ 0.025

5.3. Statistical Significance

Table 3 reports Wilcoxon signed-rank test p -values comparing the Fuzzy Ensemble against each baseline, aggregated across all five datasets. All comparisons yield $p < 0.05$ after Holm–Bonferroni correction, confirming statistical significance.

Table 3: Wilcoxon signed-rank p -values (FE vs. baselines).

Baseline	AUC	F-measure	MCC
Random Forest	0.031	0.028	0.035
XGBoost	0.024	0.021	0.027
SVM	0.008	0.006	0.009
Naive Bayes	0.003	0.002	0.004

5.4. Interpretability Analysis

Figure 5 shows Shapley-based feature importance scores derived from the learned fuzzy measure on the KC1 dataset. Lines of Code (LOC) and Cyclomatic Complexity emerge as the most influential features, consistent with established software quality research [24].


**Figure 5:** Feature importance via Fuzzy Shapley indices (KC1 dataset).

Table 4 reports classifier-level Shapley values, indicating that RF contributes most to the ensemble, followed by XGBoost and SVM, with positive synergy between RF and XGBoost ($\mu(\{RF, XGB\}) > \mu(\{RF\}) + \mu(\{XGB\})$).

Table 4: Classifier-level Shapley values and fuzzy measure subsets (KC1).

Subset	μ value	Shapley ϕ
$\{RF\}$	0.62	0.41
$\{XGB\}$	0.57	0.35
$\{SVM\}$	0.48	0.24
$\{RF, XGB\}$	0.84	—
$\{RF, SVM\}$	0.76	—
$\{XGB, SVM\}$	0.71	—
$\{RF, XGB, SVM\}$	1.00	—

5.5. Ablation Study

Table 5 presents an ablation study evaluating the impact of individual components. Removing the fuzzy measure (replacing with equal weights) reduces AUC by $\approx 2.1\%$; removing interpretability derivation has no impact on predictive metrics but eliminates explainability.

Table 5: Ablation study on KC1 (mean AUC).

Configuration	AUC
Full model (FE)	0.934
w/o fuzzy measure (equal weights)	0.913
w/o SVM base classifier	0.921
w/o XGBoost base classifier	0.916
w/o RF base classifier	0.907

6. Discussion

6.1. Performance Gains

The Fuzzy Ensemble consistently outperforms all baselines across datasets and metrics. The performance gain over the best single-model baseline (RF) averages 3.2% in AUC, 3.6% in F-measure, and 5.0% in MCC. The larger MCC gain is particularly notable given the class-imbalanced nature of SDP datasets, suggesting that the fuzzy integral’s non-linear aggregation better handles minority-class instances.

6.2. Quality of Explanations

The Shapley-based importance scores align well with domain knowledge: LOC and Cyclomatic Complexity are established defect predictors [24], validating the framework’s explanations. Unlike post-hoc LIME/SHAP applied to black-box ensembles, the fuzzy Shapley values are *intrinsic* to the model, guaranteeing consistency between the explanation and the prediction mechanism.

6.3. Practical Implications

For software quality managers, the framework provides actionable insights: (i) which modules are at risk and (ii) *why* they are flagged, through ranked feature contributions. This dual output supports prioritisation of code review and testing resources. The framework’s computational overhead over RF is modest ($<15\%$ additional training time), making it viable for industrial-scale codebases.

6.4. Limitations and Future Work

Several limitations warrant acknowledgement. First, the current implementation is limited to $n = 3$ base classifiers; scaling to larger ensembles requires approximation methods for fuzzy measure learning (e.g., k -additive measures [25]). Second, the benchmark datasets are relatively dated; evaluation on contemporary open-source projects (e.g., from GitHub) is planned. Third, the framework has not yet been validated in a controlled user study with practitioners. Future work will address these limitations and explore integration with continuous integration pipelines.

7. Conclusion

This paper presented an explainable ensemble learning framework for software defect prediction based on fuzzy-integral aggregation. By learning a fuzzy measure over base classifiers, the Choquet integral captures synergistic and redundant interactions that additive ensembles miss. The derived Shapley values provide intrinsic, model-consistent importance indices at both the classifier and feature levels.

Experiments on five NASA benchmark datasets demonstrated statistically significant improvements over RF, XGBoost, SVM, and Naive Bayes baselines in AUC, F-measure, and MCC. The qualitative analysis confirmed that the explanations align with established software quality principles, supporting practitioner trust.

The proposed framework advances the state of the art in explainable SDP and offers a principled path toward trustworthy AI-assisted software quality assurance. Source code and data will be made publicly available upon acceptance.

References

- [1] Jones C, Bonsignour O, *The Economics of Software Quality*. Addison-Wesley, 2012.
- [2] T. Hall, S. Beecham, D. Bowes, D. Gray, and S. Counsell, “A systematic literature review on fault prediction performance in software engineering,” *IEEE Trans. Software Eng.*, vol. 38, no. 6, pp. 1276–1304, 2012.
- [3] L. Breiman, “Random forests,” *Mach. Learn.*, vol. 45, no. 1, pp. 5–32, 2001.
- [4] Nalluri M, Pentela M, Eluri NR, “A scalable tree boosting system: XG boost,” in *Res. Stud. Sci. Eng. Technol*, 2020 Oct;7(12):36-51.
- [5] Chukwuere, J. E, “Exploring literature review methodologies in information systems research: A comparative study,” in *Education and Learning in Developing Nations (ELDN)*, (2023): 38-461.
- [6] Hulsén, Tim, “Explainable artificial intelligence (XAI): Explainable artificial intelligence (XAI): con-

- cepts and challenges in healthcare,” *Ai*, 4.3 (2023): 652-666.
- [7] Min, Chao, et al, “Ensemble Interpretation: A Unified Method for Interpretable Machine Learning,” in *arXiv preprint*, arXiv:2312.06255 (2023).
- [8] G. Choquet, “Theory of capacities,” *Ann. Inst. Fourier*, vol. 5, pp. 131–295, 1954.
- [9] L. S. Shapley, “A value for n -person games,” in *Contributions to the Theory of Games*, vol. 2, Princeton Univ. Press, 1953, pp. 307–317.
- [10] Menzies, Tim, “Retrospective: Data mining static code attributes to learn defect predictors,” *IEEE Transactions on Software Engineering.*, 51.3 (2025): 858-863.
- [11] Srinivasan, Kannan P., and T. Devi, “A complete and comprehensive metrics suite for object-oriented design quality assessment,” *International Journal of Software Engineering and Its Applications*, 8.2 (2014): 173-188.
- [12] M. H. Halstead, *Elements of Software Science*. Elsevier, 1977.
- [13] Priya, B. Santhi, et al, “META ANALYSIS ON CROSS PROJECT DEFECT PREDICTION,” *academia.edu.*, (2020).
- [14] Liang, Guohua, Xingquan Zhu, and Chengqi Zhang, “An empirical study of bagging predictors for different learning algorithms,” *Proceedings of the AAAI Conference on Artificial Intelligence.*, Vol. 25. No. 1. 2011.
- [15] Y. Freund and R. E. Schapire, “A decision-theoretic generalization of on-line learning and an application to boosting,” *J. Comput. Syst. Sci.*, vol. 55, no. 1, pp. 119–139, 1997.
- [16] Alrashedy, Kamel, Vincent J. Hellendoorn, and Alessandro Orso, “Learning defect prediction from unrealistic data,” in *2024 IEEE international conference on software analysis, evolution and reengineering (SANER)*, IEEE, 2024.
- [17] Abdu, Ahmed, et al, “Semantic and traditional feature fusion for software defect prediction using hybrid deep learning model,” in *Scientific Reports*, 14.1 (2024): 14771.
- [18] J. Jiarpakdee, C. Tantithamthavorn, and A. E. Hassan, “Practitioners’ perceptions of the goals and visual explanations of defect prediction models,” in *Proc. IEEE/ACM MSR*, 2021, pp. 432–443.
- [19] S. Wattanakriengkrai et al., “Predicting defective lines using a model-agnostic technique,” *IEEE Trans. Software Eng.*, vol. 48, no. 6, pp. 1–18, 2020.
- [20] M. Grabisch, “The application of fuzzy integrals in multicriteria decision making,” *Eur. J. Oper. Res.*, vol. 89, no. 3, pp. 445–456, 1996.
- [21] J. M. Keller, P. Gader, and A. Hocaoglu, “Fuzzy integrals in image processing and recognition,” in *Fuzzy Measures and Integrals*, Physica-Verlag, 2000, pp. 435–466.
- [22] D. T. Anderson et al., “Fusion of multi-spectral imagery using fuzzy integrals,” in *Proc. SPIE*, vol. 9090, 2014.
- [23] T. Menzies, B. Caglayan, E. Kocaguneli, J. Krall, F. Peters, and B. Turhan, “The PROMISE repository of empirical software engineering data,” <http://promisedata.org/>, 2012.
- [24] T. J. McCabe, “A complexity measure,” *IEEE Trans. Software Eng.*, vol. SE-2, no. 4, pp. 308–320, 1976.
- [25] M. Grabisch, “ k -order additive discrete fuzzy measures and their representation,” *Fuzzy Sets Syst.*, vol. 92, no. 2, pp. 167–189, 1997.